

# ASSIGNMENTS - 2

Assignment Date	25 September2022
Student Name	Ajay E
Team ID	PNT2022TMID25121
Maximum Marks	2 Marks

In [1]:

```
#1Load the dataset
from google.colab import drive
```

In [2]:

```
drive.mount('/content/drive')
```

Mounted at /content/drive

In [3]:

```
from pandas.io.formats.style import plt
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from matplotlib import rcParams
```

In [4]:

```
df=pd.read_csv('/content/drive/MyDrive/IBM/dataset')
```

In [5]:

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call  
drive.mount("/content/drive", force\_remount=True).

In [6]:

```
df.head()
```

Out[6]:

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	1	15634602	Hargrave	619	France	Female	42	2 0.0
1	2	15647311	Hill	608	Spain	Female	41	1 83807.8
2	3	15619304	Onio	502	France	Female	42	8 159660.8
3	4	15701354	Boni	699	France	Female	39	1 0.0
4	5	15737888	Mitchell	850	Spain	Female	43	2 125510.8

In [

7]:

```
#6 check for categorical column and perform encoding
dummy=pd.get_dummies(df['Gender'])
dummy.head()
```

Out[7]:

	Female	Male
0	1	0
1	1	0
2	1	0
3	1	0
4	1	0

In [8]:

```
df2=pd.concat((df,dummy),axis=1)
df2.head()
```

Out[8]:

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	1	Hargrave	619	France	Female	42	2	0.0
1	2	Hill	608	Spain	Female	41	1	83807.8
2	3	Onio	502	France	Female	42	8	159660.8
3	4	Boni	699	France	Female	39	1	0.0
4	5	Mitchell	850	Spain	Female	43	2	125510.8



In [

9]:

```
df2.drop(['Gender'],axis=1)
```

Out[9]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Age	Tenure	Balance	N
0	1	15634602	Hargrave	619	France	42	2	0.00	
1	2	15647311	Hill	608	Spain	41	1	83807.86	
2	3	15619304	Onio	502	France	42	8	159660.80	
3	4	15701354	Boni	699	France	39	1	0.00	
4	5	15737888	Mitchell	850	Spain	43	2	125510.82	
...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijaku	771	France	39	5	0.00	
9996	9997	15569892	Johnstone	516	France	35	10	57369.61	
9997	9998	15584532	Liu	709	France	36	7	0.00	
9998	9999	15682355	Sabbatini	772	Germany	42	3	75075.31	
9999	10000	15628319	Walker	792	France	28	4	130142.79	

10000 rows × 15 columns

In [

In [10]:

```
df2=df2.drop(['Gender'],axis=1)
df2.head()
```

Out[10]:

RowNumber	CustomerId	Surname	CreditScore	Geography	Age	Tenure	Balance	NumC
0	1	15634602	Hargrave	619	France	42	2	0.00
1	2	15647311	Hill	608	Spain	41	1	83807.86
2	3	15619304	Onio	502	France	42	8	159660.80
3	4	15701354	Boni	699	France	39	1	0.00
4	5	15737888	Mitchell	850	Spain	43	2	125510.82

In [1]:

1

```
df2=df2.drop(['Male'],axis=1)
df2.head()
```

Out[11]:

RowNumber	CustomerId	Surname	CreditScore	Geography	Age	Tenure	Balance	NumC
0	1	15634602	Hargrave	619	France	42	2	0.00
1	2	15647311	Hill	608	Spain	41	1	83807.86
2	3	15619304	Onio	502	France	42	8	159660.80
3	4	15701354	Boni	699	France	39	1	0.00
4	5	15737888	Mitchell	850	Spain	43	2	125510.82

In [1]:

In [12]:

```
df2.rename(columns={"Female":"Gender"})
```

Out[12]:

	RowNumber	CustomerId	Surname	CreditScore	Geography	Age	Tenure	Balance	N
0	1	15634602	Hargrave	619	France	42	2	0.00	
1	2	15647311	Hill	608	Spain	41	1	83807.86	
2	3	15619304	Onio	502	France	42	8	159660.80	
3	4	15701354	Boni	699	France	39	1	0.00	
4	5	15737888	Mitchell	850	Spain	43	2	125510.82	
...	...	...	...	...	...	...	...	...	...
9995	9996	15606229	Obijaku	771	France	39	5	0.00	
9996	9997	15569892	Johnstone	516	France	35	10	57369.61	
9997	9998	15584532	Liu	709	France	36	7	0.00	
9998	9999	15682355	Sabbatini	772	Germany	42	3	75075.31	
9999	10000	15628319	Walker	792	France	28	4	130142.79	

10000 rows × 14 columns

In [13]:

```
df.shape
```

Out[13]:

(10000, 14)

In [1]:

4

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   RowNumber        10000 non-null   int64  
 1   CustomerId       10000 non-null   int64  
 2   Surname          10000 non-null   object  
 3   CreditScore      10000 non-null   int64  
 4   Geography         10000 non-null   object  
 5   Gender            10000 non-null   object  
 6   Age               10000 non-null   int64  
 7   Tenure            10000 non-null   int64  
 8   Balance           10000 non-null   float64 
 9   NumOfProducts     10000 non-null   int64  
 10  HasCrCard        10000 non-null   int64  
 11  IsActiveMember    10000 non-null   int64  
 12  EstimatedSalary   10000 non-null   float64 
 13  Exited            10000 non-null   int64  
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

In [1]:

In [15]:

```
#4 Handling Missing Values
df.isnull().any()
```

Out[15]:

```
RowNumber      False
CustomerId     False
Surname        False
CreditScore    False
Geography      False
Gender         False
Age            False
Tenure          False
Balance         False
NumOfProducts   False
HasCrCard      False
IsActiveMember False
EstimatedSalary False
Exited          False
dtype: bool
```

In [1]:

6

```
df.isnull().sum()
```

Out[16]:

```
RowNumber      0
CustomerId     0
Surname        0
CreditScore    0
Geography      0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts   0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64
```

In [17]:

```
df.isnull().sum().sum()
```

Out[17]:

0

In [18]:

```
#7 split the data into independent and dependent variable
y=df2['Tenure']
y
```

Out[18]:

```
0      2
1      1
2      8
3      1
4      2
..
9995    5
9996   10
9997    7
9998    3
9999    4
Name: Tenure, Length: 10000, dtype: int64
```

In [1]:

9

```
X=df.drop(columns=['Balance'],axis=1)  
X.head()
```

Out[19]:

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	NumOfP
0	1	15634602	Hargrave	619	France	Female	42	2
1	2	15647311	Hill	608	Spain	Female	41	1
2	3	15619304	Onio	502	France	Female	42	8
3	4	15701354	Boni	699	France	Female	39	1
4	5	15737888	Mitchell	850	Spain	Female	43	2

In [1]:

In [20]:

```
X=df.iloc[:,7:10]
```

```
X
```

Out[20]:

	Tenure	Balance	NumOfProducts
0	2	0.00	1
1	1	83807.86	1
2	8	159660.80	3
3	1	0.00	2
4	2	125510.82	1
...	...	...	...
9995	5	0.00	2
9996	10	57369.61	1
9997	7	0.00	1
9998	3	75075.31	2
9999	4	130142.79	1

10000 rows × 3 columns

In [21]:

```
print(X.shape)
```

(10000, 3)

In [ ]:

22

```
Y=df.iloc[:,7]
```

```
Y
```

Out[22]:

```
0      2  
1      1  
2      8  
3      1  
4      2  
..  
9995    5  
9996   10  
9997    7  
9998    3  
9999    4  
Name: Tenure, Length: 10000, dtype: int64
```

In [23]:

```
print(y.shape)
```

(10000,)

In [24]:

```
#8 Scale the independent variable  
from sklearn.preprocessing import scale  
x_scaled=pd.DataFrame(scale(X),columns=X.columns)  
x_scaled.head()
```

Out[24]:

	Tenure	Balance	NumOfProducts
0	-1.041760	-1.225848	-0.911583
1	-1.387538	0.117350	-0.911583
2	1.032908	1.333053	2.527057
3	-1.387538	-1.225848	0.807737
4	-1.041760	0.785728	-0.911583

In [25]:

```
#9 Train test split  
from sklearn.model_selection import train_test_split
```

In [26]:

```
# Split data (train & test data)  
X_train, X_test, y_train, y_test = train_test_split(X,y)
```

In [ ]:

27

```
#display shape
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(7500, 3)
(2500, 3)
(7500,)
(2500,)
```

In [28]:

```
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.30,random_state=0)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(7000, 3)
(3000, 3)
(7000,)
(3000,)
```

In [29]:

```
print(np.mean(X_train))
print(np.mean(X_test))
```

```
Tenure           4.989429
Balance          75204.853421
NumOfProducts    1.534286
dtype: float64
Tenure           5.067333
Balance          79474.972977
NumOfProducts    1.520667
dtype: float64
```

In [ ]:

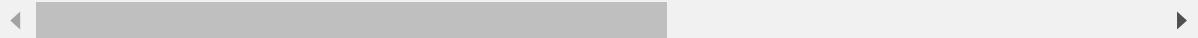
30

```
#3 Descriptive stastical analysis  
df.describe()
```

Out[30]:

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	N
<b>count</b>	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000000	
<b>mean</b>	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889288	
<b>std</b>	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405202	
<b>min</b>	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000000	
<b>25%</b>	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000000	
<b>50%</b>	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540000	
<b>75%</b>	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240000	
<b>max</b>	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090000	

In [ ]:



In [31]:

```
df.NumOfProducts.value_counts()
```

Out[31]:

```
1    5084
2    4590
3    266
4     60
Name: NumOfProducts, dtype: int64
```

In [ ]:

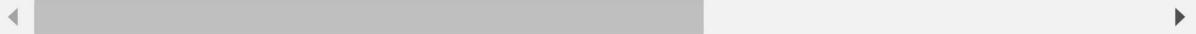
32

df.corr()

Out[32]:

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balance	NumOf
RowNumber	1.000000	0.004202	0.005840	0.000783	-0.006495	-0.009067	
CustomerId	0.004202	1.000000	0.005308	0.009497	-0.014883	-0.012419	
CreditScore	0.005840	0.005308	1.000000	-0.003965	0.000842	0.006268	
Age	0.000783	0.009497	-0.003965	1.000000	-0.009997	0.028308	-
Tenure	-0.006495	-0.014883	0.000842	-0.009997	1.000000	-0.012254	
Balance	-0.009067	-0.012419	0.006268	0.028308	-0.012254	1.000000	-
NumOfProducts	0.007246	0.016972	0.012238	-0.030680	0.013444	-0.304180	
HasCrCard	0.000599	-0.014025	-0.005458	-0.011721	0.022583	-0.014858	
IsActiveMember	0.012044	0.001665	0.025651	0.085472	-0.028362	-0.010084	
EstimatedSalary	-0.005988	0.015271	-0.001384	-0.007201	0.007784	0.012797	
Exited	-0.016571	-0.006248	-0.027094	0.285323	-0.014001	0.118533	-

In [ ]:



In [33]:

```
df['Age'].mean()
```

Out[33]:

38.9218

In [34]:

```
df['Gender'].value_counts()
```

Out[34]:

```
Male      5457  
Female    4543  
Name: Gender, dtype: int64
```

In [ ]:

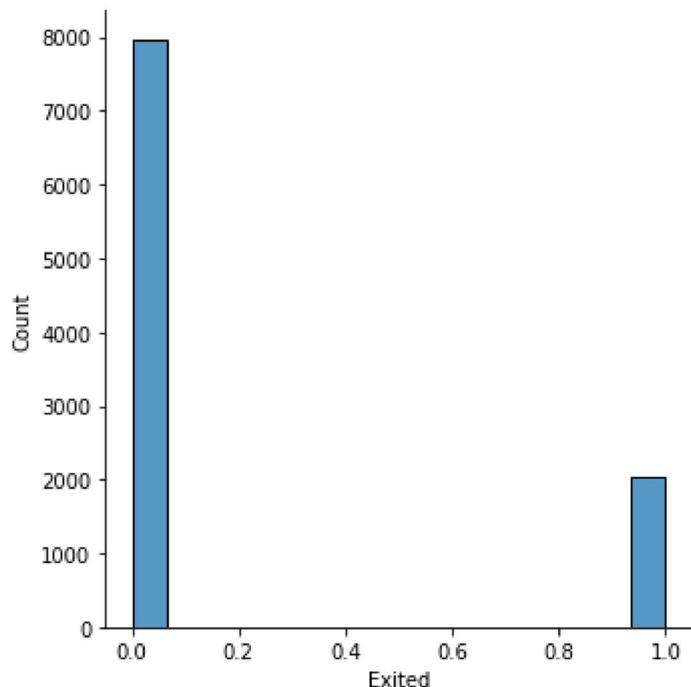
```
#2 Visualization
```

61

```
#univariate analysis  
sns.displot(df.Exited)
```

Out[61]:

```
<seaborn.axisgrid.FacetGrid at 0x7f8a972818d0>
```



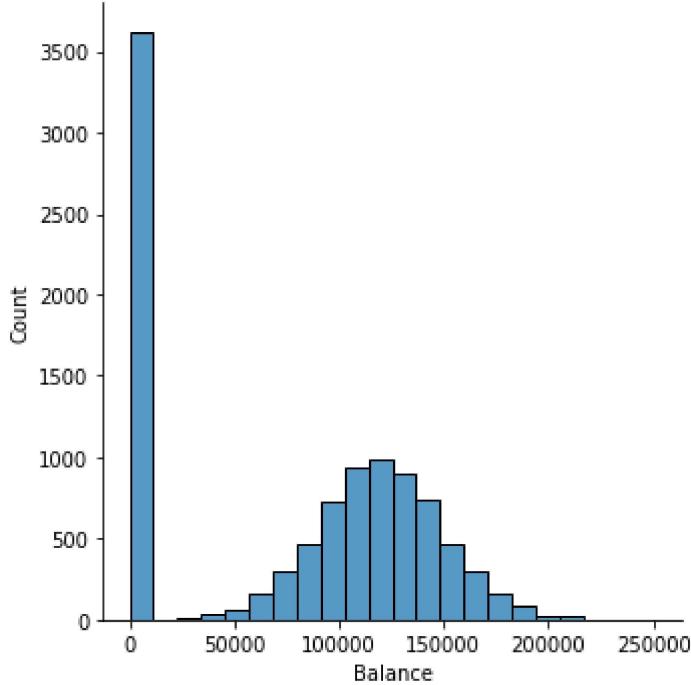
In [ ]:

60

```
sns.displot(df.Balance)
```

Out[60]:

```
<seaborn.axisgrid.FacetGrid at 0x7f8a97db4a50>
```

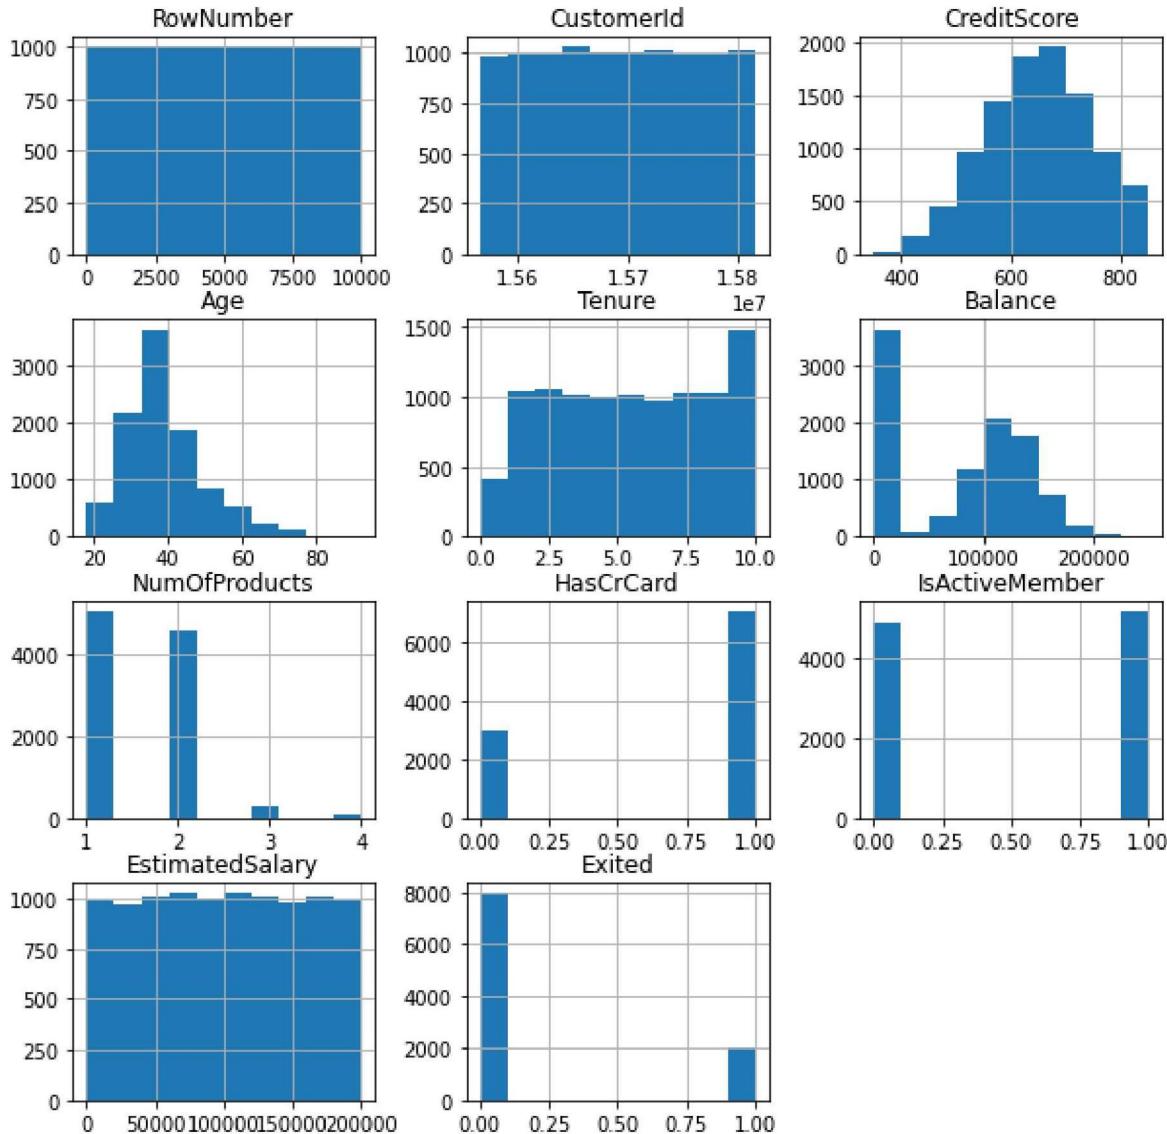


37

```
df.hist(figsize=(10,10))
```

Out[37]:

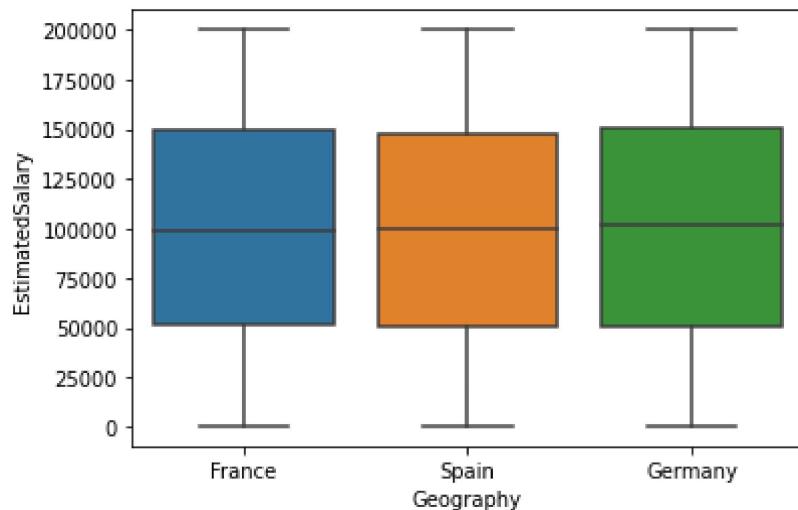
```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f8a9ca3a150>,
       <matplotlib.axes._subplots.AxesSubplot object at 0x7f8a9c9ee790>,
       <matplotlib.axes._subplots.AxesSubplot object at 0x7f8a9c9a4d90>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f8a9c96a3d0>,
       <matplotlib.axes._subplots.AxesSubplot object at 0x7f8a9c924cd0>,
       <matplotlib.axes._subplots.AxesSubplot object at 0x7f8a9c8e9550>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f8a9c8a1d10>,
       <matplotlib.axes._subplots.AxesSubplot object at 0x7f8a9c869290>,
       <matplotlib.axes._subplots.AxesSubplot object at 0x7f8a9c8692d0>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f8a9c81c550>,
       <matplotlib.axes._subplots.AxesSubplot object at 0x7f8a9c7fdb50>,
       <matplotlib.axes._subplots.AxesSubplot object at 0x7f8a9c74b690>]],
     dtype=object)
```



In [ ]:

In [62]:

```
#Boxplot  
sns.boxplot(x='Geography',y='EstimatedSalary',data=df)  
plt.show()
```



39

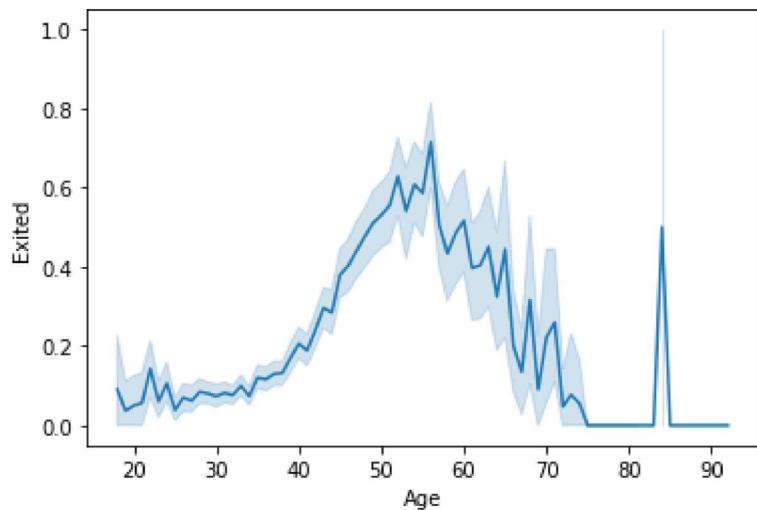
```
# Bivariate Analysis on continuous variable  
sns.lineplot(df.Age,df.Exited)
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[39]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8a9c438650>
```



In [40]:

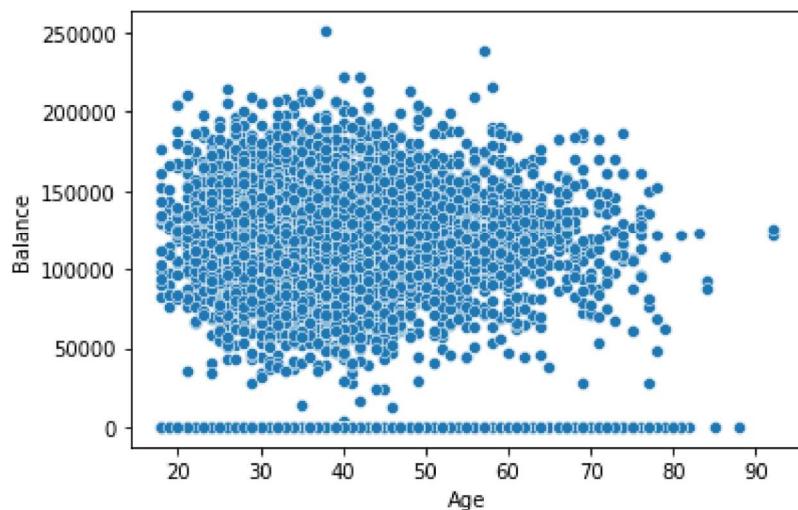
```
sns.scatterplot(df.Age,df.Balance)
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[40]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8a9c308b10>
```



In [4]:

1

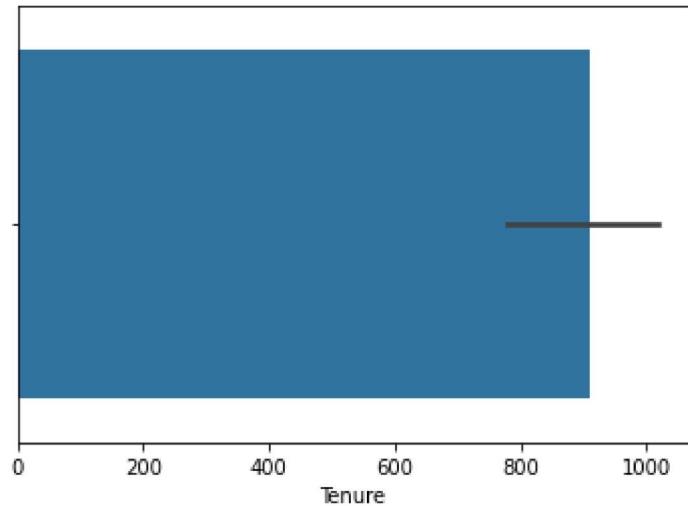
```
# Bivariate Analysis on categorical variable
sns.barplot(df.Tenure.value_counts())
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[41]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8a9c313150>
```



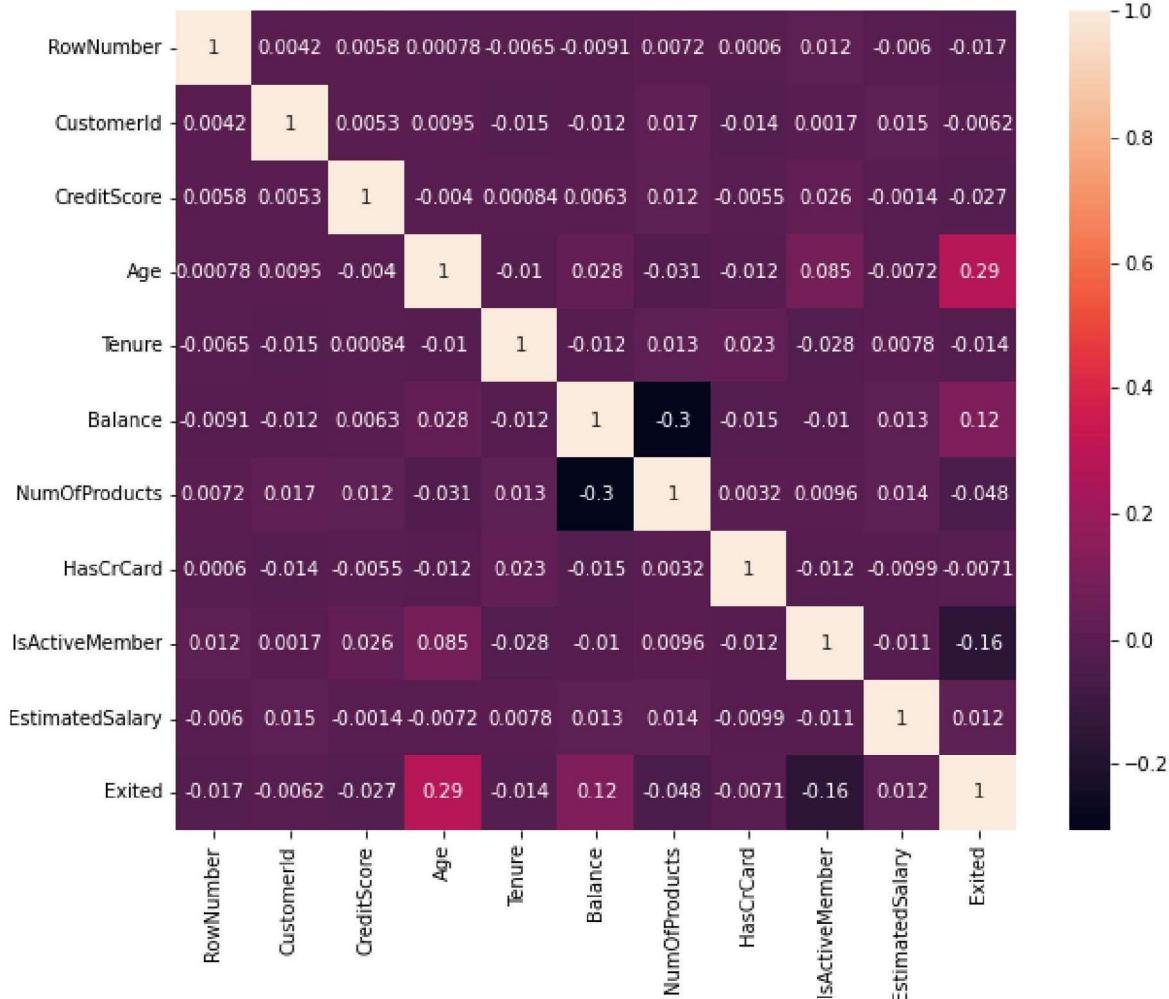
In [4] :

2

```
fig=plt.figure(figsize=(10,8))
sns.heatmap(df.corr(),annot=True)
```

Out[42]:

&lt;matplotlib.axes.\_subplots.AxesSubplot at 0x7f8a9c313050&gt;



In [4]:

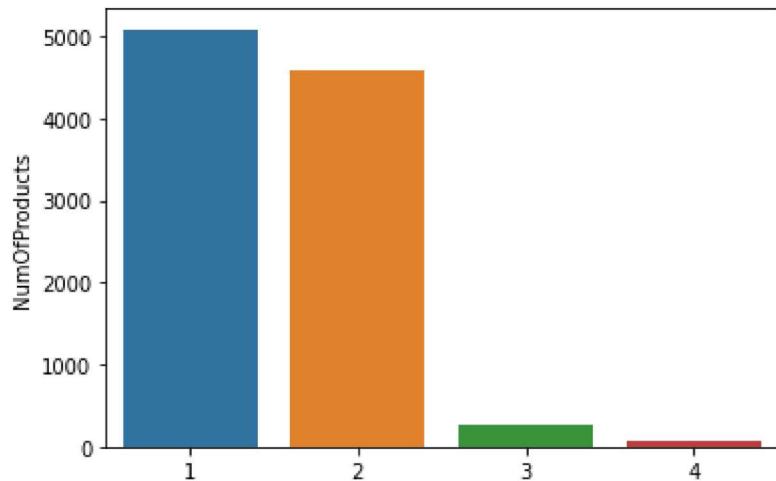
3

```
sns.barplot(df.NumOfProducts.value_counts().index,df.NumOfProducts.value_counts())
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
FutureWarning
```

Out[43]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8a9c04b6d0>
```



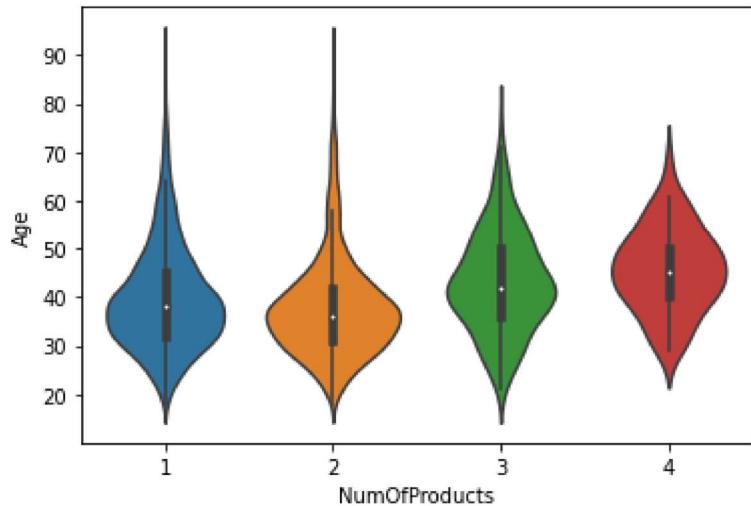
In [4]:

In [65]:

```
sns.violinplot(x='NumOfProducts',y='Age',data=df,size=8)  
plt.show
```

Out[65]:

```
<function matplotlib.pyplot.show(*args, **kw)>
```



In [4]:

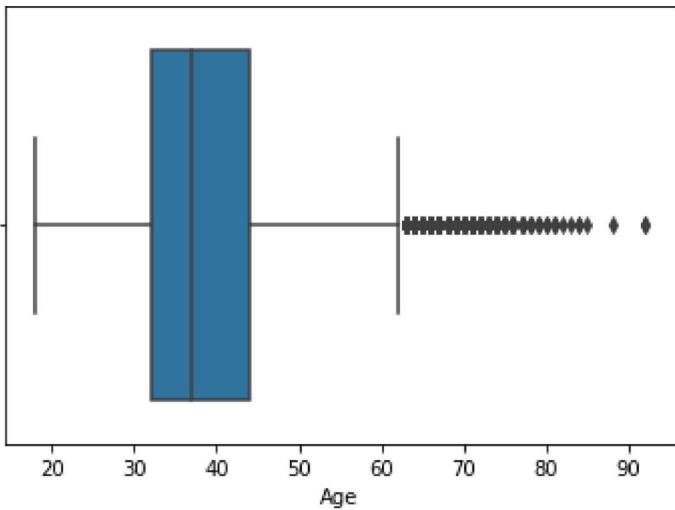
5

```
sns.boxplot(df.Age)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.  
FutureWarning
```

Out[45]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8a9a8329d0>
```

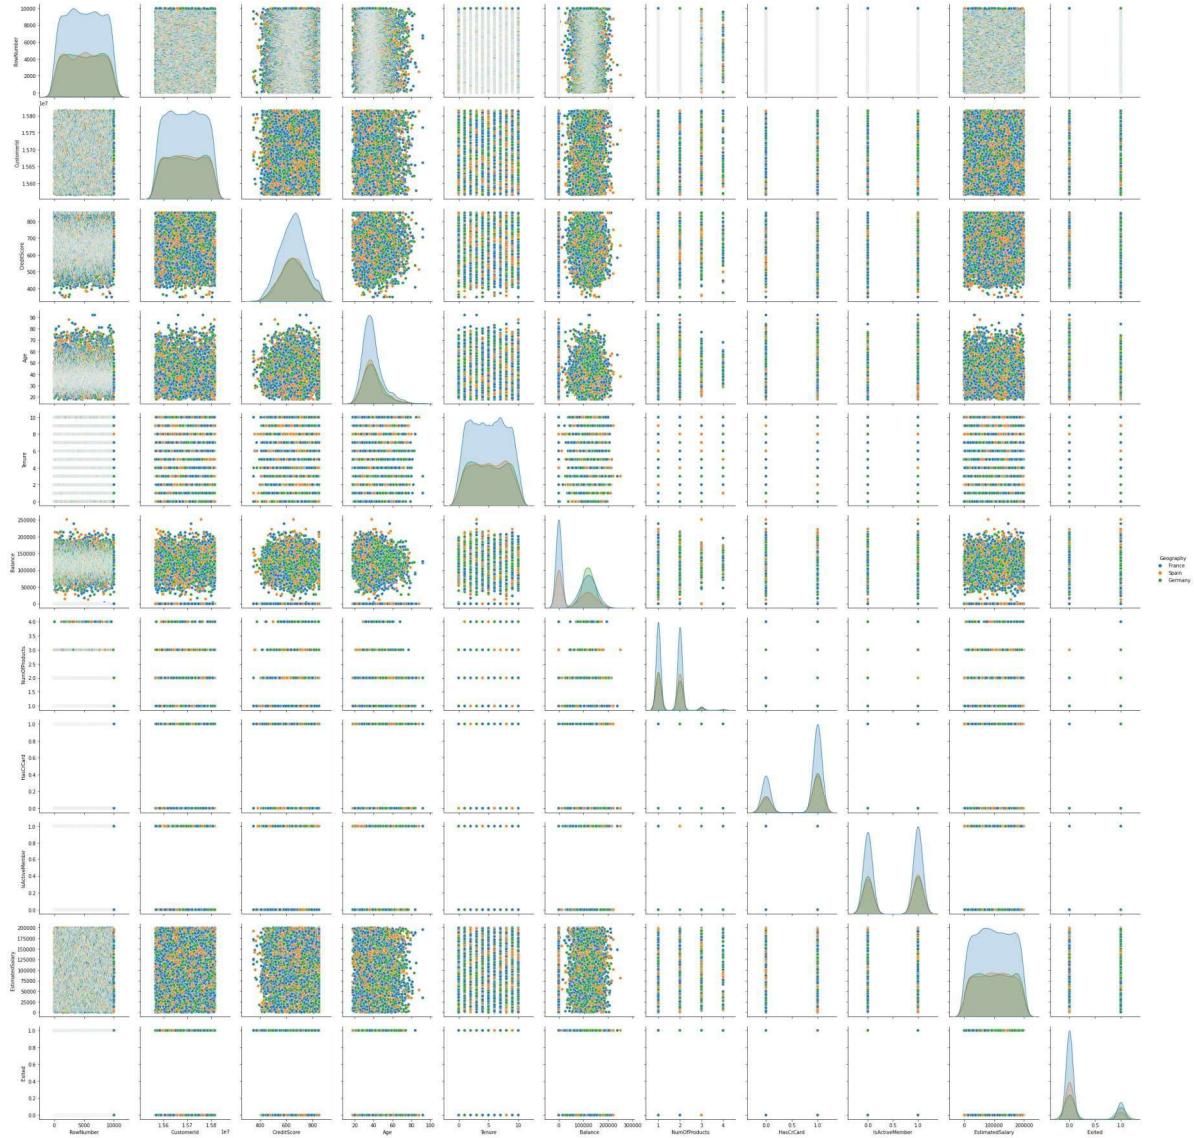


In [ ]:

66

## #Multivariate analysis

```
sns.pairplot(data=df,hue='Geography',height=3)  
plt.show()
```



In [47]:

```
#5 Find the outlier and replace the outlier
#find the limits
upper_limit=df['Age'].mean()+3*df['Age'].std()
lower_limit=df['Age'].mean()-3*df['Age'].std()
print('upper_limit',upper_limit)
print('lower_limit',lower_limit)
```

```
upper_limit 70.38521935511383
lower_limit 7.458380644886169
```

In [48]:

```
df.loc[(df['Age']>upper_limit)|(df['Age']<lower_limit)]
```

Out[48]:

RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Bal
85	86	15805254	Ndukaku	652	Spain	Female	75	10
158	159	15589975	Maclean	646	France	Female	73	6 97
230	231	15808473	Ringrose	673	France	Male	72	1
252	253	15793726	Matveyeva	681	France	Female	79	0
310	311	15712287	Pokrovskii	652	France	Female	80	4
...	...	...	...	...	...	...	...	...
9646	9647	15603111	Muir	850	Spain	Male	71	10 69
9671	9672	15636061	Pope	649	Germany	Male	78	4 68
9736	9737	15644103	Wells	659	Spain	Male	78	2 151
9894	9895	15704795	Vagin	521	France	Female	77	6
9936	9937	15653037	Parks	609	France	Male	77	1

133 rows × 14 columns

In [49]:

```
#trimming - delete the outlier
new_df=df.loc[(df['Age']<upper_limit)&(df['Age']>lower_limit)]
print('Before removing outlier:',len(df))
print('After removing outlier:',len(new_df))
```

```
Before removing outlier: 10000
After removing outlier: 9867
```

In [5]:

0

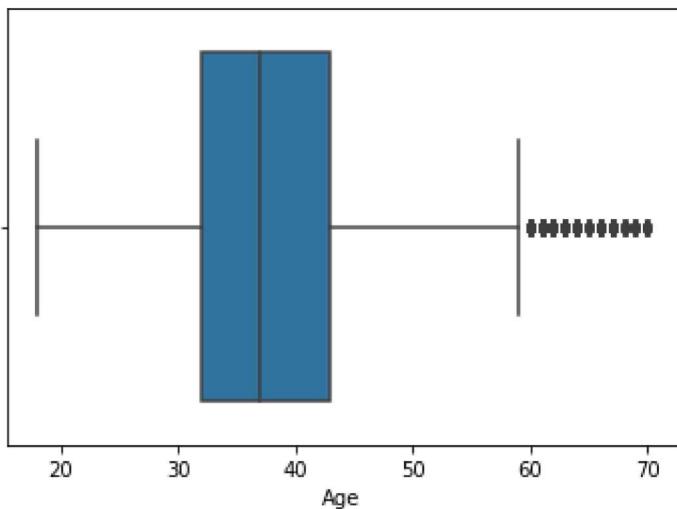
```
sns.boxplot(new_df[ 'Age' ])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```

```
FutureWarning
```

Out[50]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8a97f01790>
```



In [5]:

In [51]:

```
#copying-change the outlier value to upper or lower limits values
new_df=df.copy()
new_df.loc[(new_df['Age']>upper_limit), 'Age']=upper_limit
new_df.loc[(new_df['Age']<lower_limit), 'Age']=lower_limit
```

2

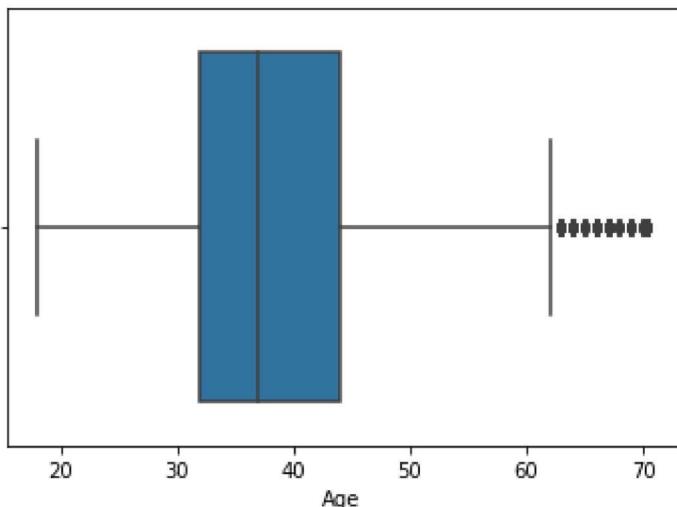
```
sns.boxplot(new_df['Age'])
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[52]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8a98031ad0>
```



In [5]:

In [53]:

```
len(new_df)
```

Out[53]:

10000

In [54]:

```
#IQR method
q1=df['Age'].quantile(0.25)
q3=df['Age'].quantile(0.75)
iqr=q3-q1
```

In [55]:

```
q1,q3,iqr
```

Out[55]:

(32.0, 44.0, 12.0)

In [5]:

6

```
upper_limit=q3+(1.5*iqr)
lower_limit=q3-(1.5*iqr)
lower_limit,upper_limit
```

Out[56]:

(26.0, 62.0)

In [57]:

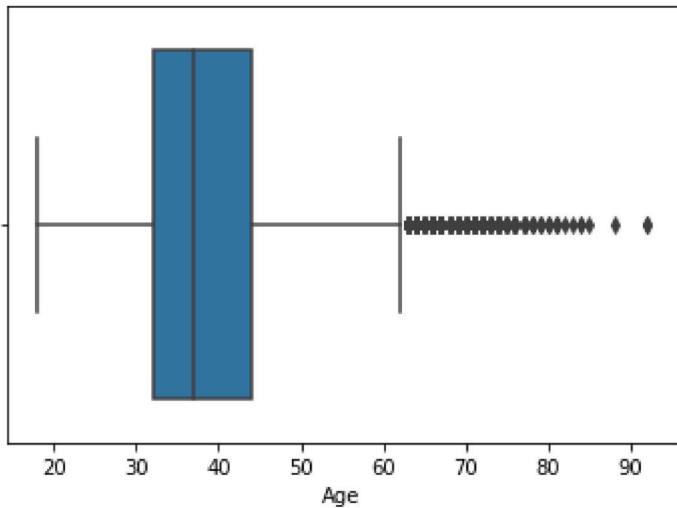
```
sns.boxplot(df['Age'])
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

Out[57]:

&lt;matplotlib.axes.\_subplots.AxesSubplot at 0x7f8a98077ed0&gt;



In [58]:

```
#copying-change the outlier value to upper or lower limits values
new_df=df.copy()
new_df.loc[(new_df['Age']>upper_limit),'Age']=upper_limit
new_df.loc[(new_df['Age']<lower_limit),'Age']=lower_limit
```

In [5]:

9

```
sns.boxplot(new_df[ 'Age' ])
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
```

```
FutureWarning
```

Out[59]:

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f8a97e1a690>
```

